

ProcTrac: Linux Kernel Module to Monitor and Track File Access

B. S. Hong

December 26, 2018

Abstract

A Linux kernel module was built to keep a list of files to keep track of and raise/log an alert when a file in that list was accessed. In this project, this was accomplished by hooking into some system calls to raise an alert.

1 Design

The design of this project is fairly simple and straightforward; on one end of the program, is an interface with the user via some means. Following the UNIX philosophy of “everything is a file,” it made sense to expose a file where the kernel module would keep its configuration (the list of files it would be monitoring for). On the other end, a handful of system calls would be hooked and whenever a syscall like *open(3)* would be called, it would check the parameters against the list and raise an alert when there was a match.

2 Challenges

Despite its simpleness, there were quite a few challenges that I’ve ran into. Most of the problem arises from writing a Linux kernel itself and often the complexity of doing so. There are many tutorials online that show how to write a basic kernel module. This was fine, but as I tried to add functionality to the module, it proved very hard for several reasons. Firstly, there weren’t as many guides or documentations on the exact subsystem or the exact function that I was looking for. To make that worse, the Linux kernel went through a major internal change on the 2.6 release. The entire linking system changed and moved from userspace to kernel space [?]. This caused half the guides to become outdated since the kernel symbols were no longer exported and even if you find the syscall table, it is also marked read-only.

Another challenge was that it was just simply too different of an environment to learn to code for, especially if not well familiar with the kernel internals.

After a bit of research, I found out about `ftrace`. `Ftrace` is a tracing utility directly built into the kernel. Using `ftrace`, I was able to install the function hooks. One of the issues that I had while coding this part is that debugging was very annoying; normally, there's a bug and your program just segfaults, terminates, or spits out an error. In the kernel space, when something goes wrong, the kernel would panic and halt. I would then have to perform a hard reset on my machine and also lose all the debugging info.

Another really hard part was the user input part. Initially, I tried to create a config file like `'/etc/ptrac.conf'`. The users would modify that file and when the kernel module gets loaded, it would read from that file. However, this proved to be harder than expected. To be fair, opening up files from the kernel is considered a really bad practice anyways, and it was rightfully difficult in attempts to dissuade bad coders to take that route. If you wish to, the commit history would show that I've written a bunch of helper functions like `open`, `write`, `read`, `close` to try to make this work, but could not get read/write to work correctly.

The next thing attempted was to expose a `sysfs` file. The `sysfs` is a virtual filesystem provided by the kernel, very much like the `/proc` or `/dev` directories in a typical Linux system, that does not map to a real filesystem, but maps to the kernel memory, somehow. It basically exposes the kernel kobject models to userspace via a virtual filesystem. This part also had very little standard documentation and the guides that were found all used different functions and stuff. Eventually, I went through the kernel docs and was able to find and use the right functions to create a `sysfs` entry that the user and the module can communicate through.

3 Implementation

As described above in the previous section, the two main parts of the project are the hooking and the `sysfs` entry. Firstly, when the module is inserted, it creates a kobject and `sysfs` directory `/sys/ptrac` (line 311). It then creates an 'attribute' of that kobject just created, pointing to the two handler functions that would read and write from the `sysfs` file `/sys/ptrac/filelist` (lines 302,314). `'__ATTR_RW()'` is a C macro included from `linux/sysfs.h`. If inspected carefully, you may notice that `filelist_store` actually reads a string for filename and an int for access. This will be further discussed in the next section.

There actually is a third part of the module which is internal, and that's the linked list that stores all the filepaths. After being read, the filename is then stored in the `struct st_fcntrl`, which is a node in a singly-linked list pointed to by `flist`. The files are added onto this list through the function `filelist_store` and removed in either `filelist_store` when the

access is 0 or when the kernel module is unloaded. The function *filelist_show* handles when a user process reads out the sysfs file; it simply traverses through the linked list and prints the access value and the full filepath.

After the sysfs entry is setup, the module installs all the hooks defined in lines 231 to 235. Additional functions can be added by adding another of these lines and creating *real_sys_name* and *hook_sys_name* like I have with other functions. They also have to be installed and removed. The functions *install_hook*, *ftrace_hook*, *register_ftrace_function*, *remove_hook*, *resolve_hook_address*, and *ftrace_thunk* are all either wrapper or helper functions that actually deal with the hooks.

In *resolve_hook_address*, there is a call to *kallsyms_lookup_name*. Since the kernel symbols are no longer statically exported, this is how the function addresses are dynamically resolved from its names.

The hooks, when called, all do very similar things: they first check if the filename provided is in the list of filenames to be monitored. If it is, then it prints an alert with the PID of the process that is trying to access the file. It then calls the original, real syscall and returns the returned value. One thing I figured out working on this part was that the decl spec *__user* actually meant something. It meant that the data was in userspace, and that it should not be dereferenced directly. That is why the call to *strncpy_from_user()* was necessary in *dup_fn*.

Things were working fine, until I tested hooking *sys_execve*. I realized that the filename would be sometimes a relative path rather than a full path. That's what the function *resolve_path* was meant to fix. It would get the full filepath based on the current working directory of the current process. The easiest method to accomplish this seemed to be to just call the syscall *sys_getcwd* after finding out that calling syscalls are allowed while in kernel space. Once again, *kallsyms_lookup_names* was used to try to resolve *sys_getcwd*, but the function just didn't work right. (I forgot to remove that line with *sys_getcwd*.) Anyways, the *resolve_path* function returns the full path in a new buffer, making sure to free all the old and intermediate buffers. (Another example of annoyance of using kernel functions: the usage for the function *d_path* was in the kernel source and had a "NOTE" which was actually pretty significant in its usage).

4 Future Possibilities

There were a few more things I wanted to add and a few more things that I could've added, but I'll list and explain them here.

The previously mentioned 'access' variable that gets read and stored along with the filename was supposed to be like the flags parameter. Different access values would keep track of different actions for that file. For example, access value of 1 may only keep track of

open() syscalls, or 2 may only keep track of removal of that file (*unlink()* and *unlinkat()*), and so on.

Something I started to deal with but didn't quite cover everything was relative paths. In *resolve_path* it removes all the leading './'s but nothing else. So anything that begins with a './' or any dot/dot-dot directories in the middle would not be handled correctly and likely will not match any filenames during the search.

Another possible feature of this module could be to actually not only monitor, but also deny certain actions. Upon looking up the filename and comparing the access values with the attempted action, the function hook can return an error value instead of calling the real syscall.

Appendix A Source Code

The full project, including the Makefile, can be found on my GitHub page
<https://github.com/awkwardbunny/proctrac>

```
1 #include <linux/init.h>
2 #include <linux/module.h>
3 #include <linux/kernel.h>
4 #include <linux/kobject.h>
5 #include <linux/slab.h>
6 #include <linux/ftrace.h>
7 #include <linux/uaccess.h>
8 #include <linux/fs_struct.h>
9
10 MODULE_LICENSE("GPL");
11 MODULE_AUTHOR("Brian Hong");
12 MODULE_DESCRIPTION("Process File Access Tracker Kernel Module");
13 MODULE_VERSION("0.1");
14
15 // Helper function to copy string from userspace
16 static char *dup_fn(const char __user *filename){
17     char *kfn;
18     kfn = kmalloc(512, GFP_KERNEL);
19     if (!kfn)
20         return NULL;
21     if (strncpy_from_user(kfn, filename, 512) < 0){
22         kfree(kfn);
23         return NULL;
24     }
25     return kfn;
26 }
```

```

27
28 static long (*sys_getcwd)(char __user *buf, unsigned long size);
29 static char *resolve_path(char *fn){
30     char *buf, *buf2;
31     char *cwd;
32     int len;
33     int offset = 0;
34
35     // No need
36     if(fn[0] == '/')
37         return fn;
38
39     while(fn[offset] == '.'){
40         if(fn[offset+1] == '.' && fn[offset+2] == '/')
41             //offset += 3;
42             // TODO: Remove top path
43             continue;
44         else if(fn[offset+1] == '/')
45             offset += 2;
46         else
47             break;
48     }
49
50     // Get cwd
51     buf = kmalloc(1024, GFP_KERNEL);
52     if(!buf) return NULL;
53     cwd = d_path(&(current->fs->pwd), buf, 1024);
54     len = strlen(cwd);
55
56     // Concatenate two halves
57     buf2 = kmalloc(1024, GFP_KERNEL);
58     strcpy(buf2, cwd); // cwd
59     buf2[len] = '/'; // '/'
60     strcpy(buf2+len+1, fn+offset); // filename
61
62     kfree(buf);
63     kfree(fn);
64     return buf2;
65 }
66
67 // Definition of data structures to keep track of files
68 typedef struct st_fcontrl fcontrl;
69 typedef struct st_fcontrl {
70     char fn[512];
71     int access;

```

```

72     fctrl *next;
73 } fctrl;
74 fctrl *flist = NULL;
75
76 // Helper function to search through file linked list
77 static int search_flist(char *filename){
78     fctrl *fcp = flist;
79     while(fcp){
80         if(!strcmp(fcp->fn, filename))
81             return fcp->access;
82         fcp = fcp->next;
83     }
84     return 0;
85 }
86
87 // Definitions and functions for hooking
88 struct ftrace_hook {
89     const char *name;
90     void *function;
91     void *original;
92
93     unsigned long address;
94     struct ftrace_ops ops;
95 };
96
97 #define HOOK(_name) \
98 { \
99     .name = #_name, \
100     .function = hook_##_name, \
101     .original = &real_##_name \
102 }
103
104 static int resolve_hook_address(struct ftrace_hook *hook){
105     hook->address = kallsyms_lookup_name(hook->name);
106     if(!hook->address){
107         printk("unresolved symbol: %s\n", hook->name);
108         return -ENOENT;
109     }
110
111     *((unsigned long *) hook->original) = hook->address;
112     return 0;
113 }
114
115 static void notrace ftrace_thunk(unsigned long ip, unsigned long parent_ip,
    struct ftrace_ops *ops, struct pt_regs *regs){

```

```

116 struct ftrace_hook *hook = container_of(ops, struct ftrace_hook, ops);
117 if(!within_module(parent_ip, THIS_MODULE))
118     regs->ip = (unsigned long) hook->function;
119 }
120
121 int install_hook (struct ftrace_hook *hook){
122     int err;
123     err = resolve_hook_address(hook);
124     if(err)
125         return err;
126     hook->ops.func = ftrace_thunk;
127     hook->ops.flags = FTRACE_OPS_FL_SAVE_REGS | FTRACE_OPS_FL_IPMODIFY;
128
129     err = ftrace_set_filter_ip(&hook->ops, hook->address, 0, 0);
130     if(err){
131         printk("ftrace_set_filter_ip() failed: %d\n", err);
132         return err;
133     }
134
135     err = register_ftrace_function(&hook->ops);
136     if(err){
137         printk("register_ftrace_function() failed: %d\n", err);
138         ftrace_set_filter_ip(&hook->ops, hook->address, 1, 0);
139         return err;
140     }
141     printk(KERN_INFO "PTRAC: Installed hook on %s()\n", hook->name);
142     return 0;
143 }
144
145 void remove_hook(struct ftrace_hook *hook){
146     int err;
147     err = unregister_ftrace_function(&hook->ops);
148     if(err)
149         printk("unregister_ftrace_function() failed: %d\n", err);
150     err = ftrace_set_filter_ip(&hook->ops, hook->address, 1, 0);
151     if(err)
152         printk("ftrace_set_filter_ip() failed: %d\n", err);
153     printk(KERN_INFO "PTRAC: Removed hook on %s()\n", hook->name);
154 }
155
156 // Hook definitions
157 static asmlinkage long (*real_sys_open)(const char __user *filename, int flags,
158     , umode_t mode);
159 static asmlinkage long hook_sys_open(const char __user *filename, int flags,
160     umode_t mode){

```

```

159     long ret;
160     char *kfn;
161
162     kfn = dup_fn(filename);
163     kfn = resolve_path(kfn);
164     if(search_flist(kfn)) printk(KERN_INFO "PTRAC: PID %d is opening %s\n",
165         task_pid_nr(current), kfn);
166     kfree(kfn);
167
168     ret = real_sys_open(filename, flags, mode);
169     return ret;
170 }
171
172 static asmlinkage long (*real_sys_unlink)(const char __user *filename);
173 static asmlinkage long hook_sys_unlink(const char __user *filename){
174     long ret;
175     char *kfn;
176
177     kfn = dup_fn(filename);
178     kfn = resolve_path(kfn);
179     if(search_flist(kfn)) printk(KERN_INFO "PTRAC: PID %d is unlinking %s\n",
180         task_pid_nr(current), kfn);
181     kfree(kfn);
182
183     ret = real_sys_unlink(filename);
184     return ret;
185 }
186
187 static asmlinkage long (*real_sys_unlinkat)(int dfd, const char __user *
188     filename, int flag);
189 static asmlinkage long hook_sys_unlinkat(int dfd, const char __user *filename,
190     int flag){
191     long ret;
192     char *kfn;
193
194     kfn = dup_fn(filename);
195     kfn = resolve_path(kfn);
196     if(search_flist(kfn)) printk(KERN_INFO "PTRAC: PID %d is unlinking %s\n",
197         task_pid_nr(current), kfn);
198     kfree(kfn);
199
200     ret = real_sys_unlinkat(dfd, filename, flag);
201     return ret;
202 }

```



```

199 static asmlinkage long (*real_sys_rename)(const char __user *filename1, const
      char __user *filename2);
200 static asmlinkage long hook_sys_rename(const char __user *filename1, const
      char __user *filename2){
201     long ret;
202     char *kfn1, *kfn2;
203
204     kfn1 = dup_fn(filename1);
205     kfn2 = dup_fn(filename2);
206     kfn1 = resolve_path(kfn1);
207     kfn2 = resolve_path(kfn2);
208     if(search_flist(kfn1) || search_flist(kfn2)) printk(KERN_INFO "PTRAC: PID %d
        is renaming %s to %s\n", task_pid_nr(current), kfn1, kfn2);
209     kfree(kfn1);
210     kfree(kfn2);
211
212     ret = real_sys_rename(filename1, filename2);
213     return ret;
214 }
215
216 static asmlinkage long (*real_sys_execve)(const char __user *filename, const
      char __user *const __user *argv, const char __user *const __user *envp);
217 static asmlinkage long hook_sys_execve(const char __user *filename, const char
      __user *const __user *argv, const char __user *const __user *envp){
218     long ret;
219     char *kfn;
220
221     kfn = dup_fn(filename);
222     kfn = resolve_path(kfn);
223     if(search_flist(kfn)) printk(KERN_INFO "PTRAC: PID %d is executing %s\n",
        task_pid_nr(current), kfn);
224     kfree(kfn);
225
226     ret = real_sys_execve(filename, argv, envp);
227     return ret;
228 }
229
230 // Hooks
231 static struct ftrace_hook open_hook = HOOK(sys_open);
232 static struct ftrace_hook unlink_hook = HOOK(sys_unlink);
233 static struct ftrace_hook unlinkat_hook = HOOK(sys_unlinkat);
234 static struct ftrace_hook rename_hook = HOOK(sys_rename);
235 static struct ftrace_hook execve_hook = HOOK(sys_execve);
236
237 // Function handlers for filelist kobject attribute

```

```

238 static ssize_t filelist_show(struct kobject *kobj, struct kobj_attribute *attr
    , char *buf){
239     fctrl *fcp = flist;
240     int r = 0;
241     int sum = 0;
242     while(fcp){
243         // printk(KERN_INFO "PTRAC: %d %s\n", fcp->access, fcp->fn);
244         r = sprintf(buf+sum, "%d %s\n", fcp->access, fcp->fn);
245         sum += r;
246         fcp = fcp->next;
247     }
248     return sum;
249 }
250
251 static ssize_t filelist_store(struct kobject *kobj, struct kobj_attribute *
    attr, const char *buf, size_t count){
252     fctrl *f;
253     fctrl *fcp = flist;
254     fctrl *prev = NULL;
255     char fn[512];
256     int access = 0;
257     fctrl *exists = NULL;
258
259     printk(KERN_INFO "PTRAC: Adding to filelist: %s", buf);
260     sscanf(buf, "%511s %d", fn, &access);
261
262     // Search if exists
263     while(fcp){
264         if(!strcmp(fn, fcp->fn)){
265             printk("File exists! ");
266             exists = fcp;
267             break;
268         }
269         prev = fcp;
270         fcp = fcp->next;
271     }
272
273     if(exists){
274         if(access){
275             //Update
276             exists->access = access;
277             printk("Updated access\n");
278         }else{
279             //Remove element
280             if(flist == exists){

```

```

281     flist = exists->next;
282 }else{
283     prev->next = exists->next;
284 }
285 kfree(exists);
286 printk("Removed file\n");
287 }
288 }else if(access){
289     // Create new fcontrl
290     f = (fcontrl *)kmallocc(sizeof(fcontrl), __GFP_FS);
291     strncpy(f->fn, fn, 511);
292     f->access = access;
293
294     // Add to list
295     f->next = flist;
296     flist = f;
297 }
298 return count;
299 }
300
301 // Kobject and stuff for filelist
302 struct kobj_attribute kattr = _ATTR_RW(filelist);
303 struct kobject *kobj_ref;
304
305 static int __init ptrac_init(void){
306
307     printk(KERN_INFO "PTRAC: Module loaded!\n");
308
309     // Setup sysfs
310     // Create new kobject and register /sys/ptrac
311     kobj_ref = kobject_create_and_add("ptrac", NULL);
312
313     // Create /sys/ptrac/filelist
314     if(sysfs_create_file(kobj_ref, &kattr.attr)){
315         printk(KERN_INFO "Cannot create sysfs file...\n");
316         return -1;
317     }
318
319     sys_getcwd = (long (*)(char __user *buf, unsigned long size))
320         kallsyms_lookup_name("sys_getcwd");
321
322     // Install hooks
323     install_hook(&open_hook);
324     install_hook(&unlink_hook);
325     install_hook(&unlinkat_hook);

```

```

325     install_hook(&rename_hook);
326     install_hook(&execve_hook);
327
328     return 0;
329 }
330
331 static void __exit ptrac_exit(void){
332
333     // Remove hooks
334     remove_hook(&open_hook);
335     remove_hook(&unlink_hook);
336     remove_hook(&unlinkat_hook);
337     remove_hook(&rename_hook);
338     remove_hook(&execve_hook);
339
340     // Decrement reference counter for /sys/ptrac
341     kobject_put(kobj_ref);
342     // Remove /sys/ptrac/filelist
343     sysfs_remove_file(kernel_kobj, &kattr.attr);
344
345     while(flist){
346         fcontrol *fcp = flist->next;
347         kfree(flist);
348         flist = fcp;
349     }
350
351     printk(KERN_INFO "PTRAC: Module unloaded!\n");
352 }
353
354 module_init(ptrac_init);
355 module_exit(ptrac_exit);

```